

WARREN D. MACEVOY

LUCK

PERSONAL NOTES

Copyright © 2015 Warren D. MacEvoy

PUBLISHED BY PERSONAL NOTES

WWW.COLORADOMESA.EDU

 [<http://creativecommons.org/licenses/by-nc-sa/4.0>]

You are free to:

- Share – copy and redistribute the material in any medium or format
- Adapt – remix, transform, and build upon the material

Under the following terms:

- Attribution – You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
- NonCommercial – You may not use the material for commercial purposes.
- ShareAlike – If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.

First printing, April 2015

Contents

List of Figures

- 1 Luck to get x heads on 8 fair coins. After the probabilities are arranged in decreasing order on the unit interval, $L(x)$ is the center point of equally probable outcomes. 11
- 2 Exact (blue) vs approximate (red) luck for normal distribution for $n = 1, 10$, and 100. 15

List of Tables

- 1 Table is arranged in increasing luck (which is decreasing probability). Getting exactly $x = 4$ heads is unlucky, requiring only $L = 14\%$ luck, while getting $x = 0$ or $x = 8$ heads is almost 100% luck. 11
- 2 Luck from two randomly generated distributions $\mu^{(x)}$ and $\mu^{(y)}$ uniformly chosen in $[0, 1]^{100}$, and $\sigma^{(x)}, \sigma^{(y)}$ are transposed squares of random 100×100 matrices. In each row, x is a sample from the $\mu^{(x)}, \sigma^{(x)}$, normal distribution, and y is from the $\mu^{(y)}, \sigma^{(y)}$ distribution. The actual values of x and y are not given, since they are very large (100 numbers each) and uninteresting. 14

Introduction

The point of these notes is to introduce an idea of “luck” that connects mathematical probability with the everyday notion.

As a motivating problem, imagine walking along the beach and asking a random person to toss a tennis ball so that it lands in the sand. The probability that it lands at some point would depend on the habits of the thrower and the details of the beach, but we can summarize this as some probability distribution, $p(x) = \rho(x)dx$, where $x \in \mathbb{R}^2$ is a suitable coordinate system for the beach in question. It would almost certainly not be a uniform distribution, and it would almost certainly not be particularly concentrated.

Traditional probability feels uncomfortable here. The chances of the ball landing at a given point is zero, and so miraculous. Yet anyone watching this process would only occasionally be surprised by the outcome.

As common (and mundane, not miraculous) such situations are, the language of statistics seems to have difficulty with the notion. Nor is it limited to continuous cases, just when there are a lot of possible outcomes. Such examples lead to non-zero but very small probabilities.

To distinguish from the more general notion of luck, note that that there is no extrinsic value on an outcome. To say something is “lucky” often means there is some value (different from the probability) associated with outcomes. However, outcomes that are the most valuable are often the least probable, and outcomes of equal probability ought to be equally lucky. This observation leads to the following definition of luck:

The luck $L(x)$ of an outcome x is the probability of getting any outcome y that is more probable than x , plus one-half the probability of getting any outcome y that is equally probable to x .

From the perspective of discussing luck, it is convenient to have few sets: $\Omega(x)$, the outcomes more likely than x , and $\omega(x)$, the outcomes equally likely to x .

These might easily exist, perhaps as another name, in the literature; I just don't know what it is called.

For a continuous probability distribution such as this, the chance of the ball landing in some small area dx near x is $p(x) = \rho(x)dx$. But the ball lands at a point, so dx is zero, so the probability $p(x) = \rho(x)dx$ is zero.

Definition 1. Omega. $\Omega(x)$ is set of outcomes more likely than x :

$$\Omega(x) = \{y \mid p(y) > p(x)\} . \quad (1)$$

We define $|\Omega(x)|$ as the probability an outcome is in $\Omega(x)$, $|\Omega(x)| = P(y \in \Omega(x))$. In the discrete case, this is

$$|\Omega(x)| = \sum_{y \in \Omega(x)} p(y), \quad (2)$$

and, in the continuous case,

$$|\Omega(x)| = \int_{\Omega(x)} \rho(y) dy . \quad (3)$$

Definition 2. omega. $\omega(x)$ is the set of outcomes equally likley to x :

$$\omega(x) = \{y \mid p(y) = p(x)\} . \quad (4)$$

Similar to $\Omega(x)$, we define $|\omega(x)|$ as the probability an outcome is in $\omega(x)$, $|\omega(x)| = P(y \in \omega(x))$.

In the discrete case, this is

$$|\omega(x)| = \sum_{y \in \omega(x)} p(y), \quad (5)$$

and, in the continuous case,

$$|\omega(x)| = \int_{\omega(x)} \rho(y) dy . \quad (6)$$

With these definitions in place, we define luck mathematically as follows:

Definition 3. Luck. The luck of an outcome is the probability getting any more likely outcome, plus half the probability of getting any equally likely outcome:

$$L(x) = |\Omega(x)| + \frac{1}{2}|\omega(x)| . \quad (7)$$

Properties of Luck.

- $0 \leq L(x) \leq 1$. This ranges from no luck to perfect luck.
- If $L(x)$ is close to 1, then $p(x)$ is comparatively small, and most outcomes would have a higher probability (you are lucky).
- If $L(x)$ is close to 0, then $p(x)$ is comparatively large, and most outcomes would have a lower probability (you are unlucky).
- $E(L) = \frac{1}{2}$. On average, luck is always 50:50.

For the typical case of many outcomes with different probabilities, $|\omega(x)|$ is small. For example, $|\omega(x)| = 0$ for any multivariate normal distribution.

We are interested in cases which have many possible outcomes with low but somewhat different probabilities (like the tennis ball on the beach). If the space is well divided (so $\max |\omega| = \max_x |\omega(x)|$ is small), then there are other interesting properties of luck:

- $E(f(L)) = \int_0^1 f(L) dL + \varepsilon$, where $|\varepsilon| \leq \max |f''| \cdot \max |\omega|^2 / 12$.
- For $p > 1$, $E(L^p) = 1/(p+1) - \varepsilon$, $0 \leq \varepsilon \leq p^2 \max |\omega|^2 / 12$.
- For $0 \leq a \leq b \leq 1$, $E(L \in [a, b]) = b - a + \varepsilon$, $|\varepsilon| \leq \max |\omega|$.

Example 1. Coins. Suppose we toss 8 fair coins. The probability of getting exactly x heads out of 8 tosses is given by the binomial distribution

$$p(x) = \frac{8!}{x!(8-x)!} \left(\frac{1}{2}\right)^8 \quad (8)$$

What is the luck associated with this distribution?

x	$p(x)$	$L(x)$
4	0.2734	0.1367
3 or 5	0.2188	0.4922
2 or 6	0.1094	0.8203
1 or 7	0.0313	0.9609
0 or 8	0.0039	0.9961

In particular $|\omega(x)| = 0$ for the normal, exponential, beta, and gamma distributions. As a worst-case counterexample, the flattest distribution is the uniform distribution, for which $|\omega(x)| = 1$.

Table 1: Table is arranged in increasing luck (which is decreasing probability). Getting exactly $x = 4$ heads is unlucky, requiring only $L = 14\%$ luck, while getting $x = 0$ or $x = 8$ heads is almost 100% luck.




Figure 1: Luck to get x heads on 8 fair coins. After the probabilities are arranged in decreasing order on the unit interval, $L(x)$ is the center point of equally probable outcomes.

Normal Distribution

Suppose we are in a probability space well approximated by the multivariate normal (Gaussian) distribution of a random variable $x \in \mathbb{R}^n$ with mean μ and non-singular covariance σ :

$$P_{\text{normal}}(x; \mu, \sigma) = \frac{e^{-\frac{1}{2}(x-\mu)^T \sigma^{-1}(x-\mu)}}{\sqrt{(2\pi)^n \det \sigma}}, \quad (9)$$

where

$$\mu_i = E(x_i), \quad (10)$$

and

$$\sigma_{ij} = E((x_i - \mu_i)(x_j - \mu_j)). \quad (11)$$

Suppose an observation x is made, how lucky is it?

$$L(x) = \int_{\Omega(x)} P_{\text{normal}}(y; \mu, \sigma) dy \quad (12)$$

where

$$\Omega(x) = \{y \in \mathbb{R}^n | P_{\text{normal}}(y) > P_{\text{normal}}(x)\} \quad (13)$$

$$= \left\{y \mid \left| \sqrt{\sigma^{-1}}(y - \mu) \right| < \left| \sqrt{\sigma^{-1}}(x - \mu) \right| \right\} \quad (14)$$

Note, for any $x \in \mathbb{R}^n$, $\omega(x)$ has zero volume, and so $|\omega(x)| = 0$.

So now let's compute $L(x)$, starting from the definition:

$$L(x) = |\Omega(x)| + \frac{1}{2} |\omega(x)| \quad (15)$$

$$= |\Omega(x)| + 0 \quad (16)$$

$$= \int_{\Omega(x)} P_{\text{normal}}(y; \mu, \sigma) dy \quad (17)$$

changing variables to $z = \sqrt{\sigma^{-1}}(x - \mu)$,

$$= \int_{|w| < |z|} P_{\text{normal}}(w; 0, I) dw \quad (18)$$

converting to spherical coordinates with $r_0 = |\sqrt{\sigma^{-1}}(x - \mu)|$,

$$= \frac{1}{\sqrt{(2\pi)^n}} \int_0^{r_0} \frac{n\pi^{n/2}}{\Gamma(\frac{n}{2} + 1)} r^{n-1} e^{-\frac{1}{2}r^2} dr \quad (19)$$

$$= \frac{\gamma(n/2, r_0^2/2)}{\Gamma(n/2)} \quad (20)$$

The last form uses the lower incomplete gamma function, defined to be

$$\gamma(s, x) = \int_0^x t^{s-1} e^{-t} dt \quad (21)$$

For any value of n , but particularly for large values, we find the following approximation to be very good:

$$L(x) \approx \frac{1}{2} \left[1 + \operatorname{erf}(r(x) - \sqrt{n-1/2}) \right], \quad (22)$$

where

$$r(x) = |\sqrt{\sigma^{-1}}(x - \mu)| \quad (23)$$

Not only is this result pretty, it is very useful. Suppose we have distribution parameters μ and σ , and would like to know if they fit actual observations. A traditional approach requires a large sample to estimate μ and σ , but we don't need this, nor do we need to assume that the distribution is normal. We just need to ask if the observations are surprising (lucky or unlucky). In large dimensions, numerical experiments suggest one sample is in most cases sufficient to establish practical certainty (probability of error less than 10^{-15}).

This comes from a taylor expansion of the log of the integrand in (19)

$L^{(x)}(x)$	$L^{(y)}(x)$	$L^{(x)}(y)$	$L^{(y)}(y)$
0.501 417 202 0	1.000 000 000 0	1.000 000 000 0	0.838 180 264 1
0.731 421 266 5	1.000 000 000 0	1.000 000 000 0	0.239 258 143 2
0.982 563 033 9	1.000 000 000 0	1.000 000 000 0	0.271 695 512 7
0.033 455 080 7	1.000 000 000 0	1.000 000 000 0	0.421 320 625 9
0.689 429 934 0	1.000 000 000 0	1.000 000 000 0	0.074 461 655 7
0.736 397 593 7	1.000 000 000 0	1.000 000 000 0	0.294 050 728 4
0.304 521 296 7	1.000 000 000 0	1.000 000 000 0	0.707 849 014 7
0.231 111 574 4	1.000 000 000 0	1.000 000 000 0	0.290 313 093 2
0.585 247 719 9	1.000 000 000 0	1.000 000 000 0	0.636 902 202 8
0.214 552 926 1	1.000 000 000 0	1.000 000 000 0	0.268 989 787 4

Table 2: Luck from two randomly generated distributions $\mu^{(x)}$ and $\mu^{(y)}$ uniformly chosen in $[0, 1]^{100}$, and $\sigma^{(x)}, \sigma^{(y)}$ are transposed squares of random 100×100 matrices. In each row, x is a sample from the $\mu^{(x)}, \sigma^{(x)}$, normal distribution, and y is from the $\mu^{(y)}, \sigma^{(y)}$ distribution. The actual values of x and y are not given, since they are very large (100 numbers each) and uninteresting.




Figure 2: Exact (blue) vs approximate (red) luck for normal distribution for $n = 1, 10$, and 100 .

Computation

Unfortunately, the explicit computations for luck may be impractical. However, if n independent samples are taken $S = (x_1, \dots, x_n)$, then it can be estimated as:

$$\begin{aligned}\ell(x_0) = & \frac{1}{n} \{\# \text{ of outcomes in } S \text{ more probable than } x\} \\ & + \frac{1}{2n} \{\# \text{ of outcomes in } S \text{ equally probable to } x\}\end{aligned}$$

It is a reasonably straightforward calculation to show that

$$E(\ell) = L, \tag{24}$$

and

$$E((\ell - L)^2) \leq \frac{1}{n} L \cdot (1 - L). \tag{25}$$

